

```
In [2]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [3]: df = pd.read_csv('yield_df.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Unnamed: 0	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
0	0	Albania	Maize	1990	36613	1485.0	121.0	16.37
1	1	Albania	Potatoes	1990	66667	1485.0	121.0	16.37
2	2	Albania	Rice, paddy	1990	23333	1485.0	121.0	16.37
3	3	Albania	Sorghum	1990	12500	1485.0	121.0	16.37
4	4	Albania	Soybeans	1990	7000	1485.0	121.0	16.37

```
In [5]: df.drop('Unnamed: 0',axis=1,inplace=True)
```

```
In [6]: df.shape
```

```
Out[6]: (28242, 7)
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28242 entries, 0 to 28241
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Area                                  28242 non-null  object
1   Item                                  28242 non-null  object
2   Year                                  28242 non-null  int64
3   hg/ha_yield                          28242 non-null  int64
4   average_rain_fall_mm_per_year        28242 non-null  float64
5   pesticides_tonnes                    28242 non-null  float64
6   avg_temp                             28242 non-null  float64
dtypes: float64(3), int64(2), object(2)
memory usage: 1.5+ MB
```

```
In [8]: df.isnull().sum()
```

```
Out[8]: Area                                0
Item                                          0
Year                                          0
hg/ha_yield                                0
average_rain_fall_mm_per_year              0
pesticides_tonnes                          0
avg_temp                                   0
dtype: int64
```

```
In [9]: df.duplicated().sum()
```

```
Out[9]: 2310
```

```
In [10]: df.drop_duplicates(inplace=True)
```

```
In [11]: df.duplicated().sum()
```

```
Out[11]: 0
```

Transforming average_rain_fall_mm_per_year

In summary, this code identifies the indices of rows in the DataFrame `df` where the values in the column 'average_rain_fall_mm_per_year' are not numeric strings. These rows can be considered for removal or further processing, depending on the specific use case.

```
In [12]: def isStr(obj):
          try:
              float(obj)
              return False
          except:
              return True
          to_drop = df[df['average_rain_fall_mm_per_year'].apply(isStr)].index
```

```
In [13]: df = df.drop(to_drop)
```

```
In [14]: df
```

```
Out[14]:
```

	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
0	Albania	Maize	1990	36613	1485.0	121.00	16.37
1	Albania	Potatoes	1990	66667	1485.0	121.00	16.37
2	Albania	Rice, paddy	1990	23333	1485.0	121.00	16.37
3	Albania	Sorghum	1990	12500	1485.0	121.00	16.37
4	Albania	Soybeans	1990	7000	1485.0	121.00	16.37
...
28237	Zimbabwe	Rice, paddy	2013	22581	657.0	2550.07	19.76
28238	Zimbabwe	Sorghum	2013	3066	657.0	2550.07	19.76
28239	Zimbabwe	Soybeans	2013	13142	657.0	2550.07	19.76
28240	Zimbabwe	Sweet potatoes	2013	22222	657.0	2550.07	19.76
28241	Zimbabwe	Wheat	2013	22888	657.0	2550.07	19.76

25932 rows × 7 columns


```
In [15]: df['average_rain_fall_mm_per_year'] = df['average_rain_fall_mm_per_year'].astype(np.float64)
```

Graph Frequency vs Area

```
In [16]: len(df['Area'].unique())
```

```
Out[16]: 101
```

```
In [17]: plt.figure(figsize=(15,20))  
sns.countplot(y=df['Area'])  
plt.show()
```



```
In [18]: (df['Area'].value_counts() < 500).sum()
```

```
Out[18]: 91
```

yield_per_country

```
In [19]: country = df['Area'].unique()
yield_per_country = []
for state in country:
    yield_per_country.append(df[df['Area']==state]['hg/ha_yield'].sum())
```

```
In [20]: df['hg/ha_yield'].sum()
```

```
Out[20]: 1996196943
```


```
In [21]: yield_per_country
```

```
Out[21]: [5711536,
6711464,
5722563,
32864032,
4524100,
109111062,
10852258,
4608380,
4384717,
4443889,
7720159,
4704812,
8442270,
470651,
136340329,
6263075,
6083337,
7031146,
10717883,
34706922,
4255627,
20561214,
26927138,
6083167,
10701651,
25312166,
21315591,
36828848,
16855944,
1452416,
3595638,
6210668,
15790618,
38780463,
9260371,
14571925,
16508723,
7975743,
8361103,
8619194,
10920131,
8824110,
274219558,
68067328,
10984722,
9104030,
32280700,
19698007,
100924145,
9965212,
14391737,
3698588,
6956804,
4258932,
9016288,
4174348,
6103523,
8346715]
```

Yield Per Country Graph

```
In [22]: plt.figure(figsize=(15, 20))
sns.barplot(y=country, x=yield_per_country)
```

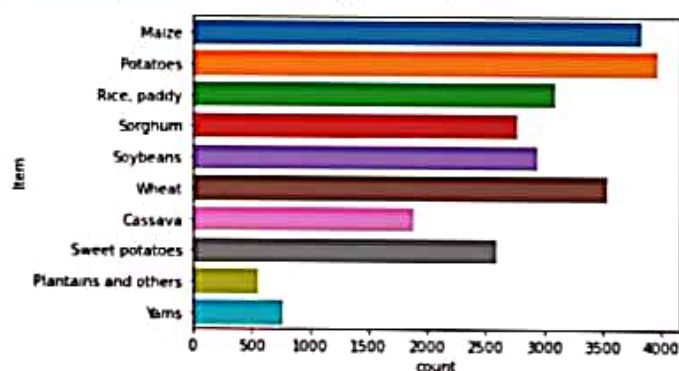
Out[22]: <AxesSubplot:>



Graph Frequency vs Item

```
In [23]: sns.countplot(y=df['Item'])
```

Out[23]: <AxesSubplot: xlabel='count', ylabel='Item'>




Yield Vs Item

```
In [24]: crops = df['Item'].unique()
yield_per_crop = []
for crop in crops:
    yield_per_crop.append(df[df['Item']==crop]['hg/ha_yield'].sum())
```

```
In [25]: sns.barplot(y=crops,x=yield_per_crop)
```

Out[25]: <AxesSubplot:~>



Train Test split Rearranging Columns

```
In [26]: col = ['Year', 'average_rain_fall_mm_per_year', 'pesticides_tonnes', 'avg_temp', 'Area', 'Item', 'hg/ha_']
df = df[col]
X = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

```
In [27]: df.head(3)
```

```
Out[27]:
```

	Year	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp	Area	Item	hg/ha_yield
0	1990	1485.0	121.0	16.37	Albania	Maize	36613
1	1990	1485.0	121.0	16.37	Albania	Potatoes	66667
2	1990	1485.0	121.0	16.37	Albania	Rice, paddy	23333

```
In [28]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8, random_state=0, shuffle=True)
```

Converting Categorical to Numerical and Scaling the values

```
In [29]: from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler
ohe = OneHotEncoder(drop='first')
scale = StandardScaler()

preprocessor = ColumnTransformer(
    transformers = [
        ('StandardScale', scale, [0, 1, 2, 3]),
        ('OHE', ohe, [4, 5]),
    ],
    remainder='passthrough'
)
```

```
In [30]: X_train_dummy = preprocessor.fit_transform(X_train)
X_test_dummy = preprocessor.transform(X_test)
```

```
In [31]: preprocessor.get_feature_names_out(col[:-1])
```

```
Out[31]: array(['StandardScale__Year',
'StandardScale__average_rain_fall_mm_per_year',
'StandardScale__pesticides_tonnes', 'StandardScale__avg_temp',
'OHE__Area_Algeria', 'OHE__Area_Angola', 'OHE__Area_Argentina',
'OHE__Area_Armenia', 'OHE__Area_Australia', 'OHE__Area_Austria',
'OHE__Area_Azerbaijan', 'OHE__Area_Bahamas', 'OHE__Area_Bahrain',
'OHE__Area_Bangladesh', 'OHE__Area_Belarus', 'OHE__Area_Belgium',
'OHE__Area_Botswana', 'OHE__Area_Brazil', 'OHE__Area_Bulgaria',
'OHE__Area_Burkina Faso', 'OHE__Area_Burundi',
'OHE__Area_Cameroon', 'OHE__Area_Canada',
'OHE__Area_Central African Republic', 'OHE__Area_Chile',
'OHE__Area_Colombia', 'OHE__Area_Croatia', 'OHE__Area_Denmark',
'OHE__Area_Dominican Republic', 'OHE__Area_Ecuador',
'OHE__Area_Egypt', 'OHE__Area_El Salvador', 'OHE__Area_Eritrea',
'OHE__Area_Estonia', 'OHE__Area_Finland', 'OHE__Area_France',
'OHE__Area_Germany', 'OHE__Area_Ghana', 'OHE__Area_Greece',
'OHE__Area_Guatemala', 'OHE__Area_Guinea', 'OHE__Area_Guyana',
'OHE__Area_Haiti', 'OHE__Area_Honduras', 'OHE__Area_Hungary',
'OHE__Area_India', 'OHE__Area_Indonesia', 'OHE__Area_Iraq',
'OHE__Area_Ireland', 'OHE__Area_Italy', 'OHE__Area_Jamaica',
```

Let's train our model

```
In [32]: #linear regression
from sklearn.linear_model import LinearRegression,Lasso,Ridge
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_absolute_error,r2_score

models = {
    'lr':LinearRegression(),
    'lss':Lasso(),
    'Rid':Ridge(),
    'Dtr':DecisionTreeRegressor()
}
for name, md in models.items():
    md.fit(X_train_dummy,y_train)
    y_pred = md.predict(X_test_dummy)

    print(f'{name} : mae : {mean_absolute_error(y_test,y_pred)} score : {r2_score(y_test,y_pred)}')

lr : mae : 29907.53512614917 score : 0.7473117803683427
C:\Users\naimat\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:592: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 67280771030.02734, tolerance: 14848622817.505226
  model = cd_fast.sparse_enet_coordinate_descent(
lss : mae : 29093.99762450549 score : 0.7473261756207235
Rid : mae : 29064.88375663324 score : 0.7473044447803092
Dtr : mae : 3817.598419124735 score : 0.9808630051556833
```

Select model

```
In [33]: dtr = DecisionTreeRegressor()
dtr.fit(X_train_dummy,y_train)
dtr.predict(X_test_dummy)

Out[33]: array([35286., 22814., 19295., ..., 16135., 34879., 77391.] )
```

Predictive System

```
In [34]: def prediction(Year, average_rain_fall_mm_per_year, pesticides_tonnes, avg_temp, Area, Item):
# Create an array of the input features
features = np.array([[Year, average_rain_fall_mm_per_year, pesticides_tonnes, avg_temp, Area, Item]])

# Transform the features using the preprocessor
transformed_features = preprocessor.transform(features)

# Make the prediction
predicted_yield = dtr.predict(transformed_features).reshape(1, -1)

return predicted_yield[0]

Year = 1990
average_rain_fall_mm_per_year = 1485.0
pesticides_tonnes = 121.00
avg_temp = 16.37
Area = 'Albania'
Item = 'Maize'
result = prediction(Year, average_rain_fall_mm_per_year, pesticides_tonnes, avg_temp, Area, Item)

C:\Users\naimat\anaconda3\lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
C:\Users\naimat\anaconda3\lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but OneHotEncoder was fitted with feature names
  warnings.warn(
```

```
In [35]: result

Out[35]: array([36613.])

In [36]: 1990    1485.0  121.00  16.37  Albania Maize    36613
2013     657.0   2550.07  19.76   Zimbabwe          Sorghum 3066
```

```
Input In [36]
1990    1485.0  121.00  16.37  Albania Maize    36613
A
SyntaxError: invalid syntax
```

Pickle Files

```
In [37]: import pickle
pickle.dump(dtr,open('dtr.pkl','wb'))
pickle.dump(preprocessor,open('preprocessor.pkl','wb'))
```

```
In [38]: import sklearn
print(sklearn.__version__)
```

1.2.2

In []: